

## **Illusory and spurious correlations: Distinct phenomena or joint outcomes of exemplar-based category learning?**

THORSTEN MEISER<sup>1\*</sup> AND MILES HEWSTONE<sup>2</sup>

<sup>1</sup>*University of Jena, Germany*

<sup>2</sup>*University of Oxford, UK*

### *Abstract*

*Stereotype formation about novel groups was analyzed with trivariate stimulus distributions that were generated by group membership, valence of behavior, and a context variable. Within this stimulus setting, we manipulated the confounding role of the context variable and the distinctiveness of events in terms of their relative infrequency. The experimental procedure allowed us to analyze illusory and spurious correlations in a joint framework, to conduct focused tests for memory effects of relative infrequency and to investigate the detection of covariations with the context variable. The results revealed that illusory and spurious correlations were formed without enhanced memory for infrequent events and with existing covariations of the confounding context factor being well extracted. These observations suggest that illusory and spurious correlations can be understood without assuming specific cognitive processes that are tied to the particular characteristics of a given stimulus distribution, such as enhanced memory in the case of relative infrequency and neglect of a context variable in the case of a confounding factor. Instead, computer simulations with an exemplar-based learning model demonstrated that exemplar-based category learning may provide a coherent and integrative theoretical framework for illusory correlations, spurious correlations and true contingency learning in social cognition. Copyright © 2006 John Wiley & Sons, Ltd.*

The cognitive processes that are involved in the abstraction of biased group stereotypes from information about individual group members have largely been investigated in two experimental paradigms: the paradigm of distinctiveness-based illusory correlations (Hamilton & Gifford, 1976) and the paradigm of spurious correlations (Schaller & O'Brien, 1992). In both paradigms, biased group impressions arise on the basis of a series of positive and negative events concerning exemplars of two novel groups. In the case of illusory correlations, an actual zero correlation between group membership and valence of behavior is misjudged because both variables are skewed. In the case of

\*Correspondence to: Thorsten Meiser, Department of Psychology, University of Jena, Humboldtstr. 11, D-07743 Jena, Germany.  
E-mail: thorsten.meiser@uni-jena.de.

Contract/grant sponsor: Alexander von Humboldt Foundation.

spurious correlations, a confounding context factor induces perception of a correlation that contradicts the actual relation between group membership and valence within each context.

Despite their procedural similarities, the two paradigms have thus far been treated separately, and distinct explanations have been proposed for the empirical phenomena of illusory and spurious correlations. The goal of the present research, in contrast, is to investigate the two phenomena in a joint approach. In a first step, we report an experiment that tested assumptions about the cognitive processes in the formation of illusory and spurious correlations that are based on specific characteristics of the stimulus distributions used in either paradigm. In a second step, we demonstrate that an exemplar-based learning account may accommodate erroneous stereotype formation across the different paradigms as a unifying theoretical framework.

## ILLUSORY CORRELATIONS

In the paradigm of distinctiveness-based illusory correlations (Hamilton & Gifford, 1976), desirable and undesirable behavior statements are presented which refer to members of two artificial groups labeled 'Group A' and 'Group B.' The overall ratio of statements about the target groups A and B as well as the overall ratio of desirable and undesirable behaviors is about 2:1. Importantly, the proportions of desirable and undesirable behaviors are constant across the two groups, yielding complete independence of group membership and desirability. Despite the independence in the stimulus set, participants perceive an illusory correlation inasmuch as they judge Group A, to which the majority of statements pertain, more favorably than Group B, to which the minority of statements pertain (see Mullen & Johnson, 1990, for a meta-analytic review). The different judgments of the Groups A and B are indicated by trait ratings, behavior assignments, and frequency estimates.

The original explanation of illusory correlations rests on the relative infrequency of Group B and of undesirable behaviors. It was argued that relative infrequency increases the salience of events, and that the co-occurrence of infrequent events attracts particular attention. Because statements that relate members of Group B to undesirable behaviors are 'paired infrequent', or 'paired distinctive', they should be more salient than other combinations of group membership and (un-)desirability (Hamilton & Gifford, 1976). As a consequence, these paired distinctive events should come to mind more easily than other events, resulting in an overestimation of the co-occurrence of Group B and undesirable behaviors in frequency judgments and group evaluations (Hamilton, 1981; Jones, Scott, Solernou, Noble, Fiala, & Miller, 1977). The distinctiveness account of illusory correlations gained support from studies that revealed extended processing of paired infrequent information (Stroessner, Hamilton, & Mackie, 1992), a recall advantage of infrequent events (Hamilton, Dugan, & Trolier, 1985; McConnell, Sherman, & Hamilton, 1994), and faster assignments of undesirable behaviors to the minority group (Johnson & Mullen, 1994; McConnell et al., 1994).

Although several studies corroborated the distinctiveness account of illusory correlations, some of the results have since been challenged. In particular, model-based analyses of recognition memory that separated memory accuracy and stereotypic guessing processes did not yield support for the presumed memory advantage for statements about the infrequent Group B or for paired distinctive events (Fiedler, Russer, & Gramm, 1993; Klauer & Meiser, 2000; Meiser & Hewstone, 2001). These analyses could only replicate the memory advantage for infrequent undesirable behaviors. Because undesirability and relative infrequency were confounded in these studies as in the previous experiments on illusory correlations, however, one cannot decide whether the observed memory advantage for undesirable behaviors was due to the distinctiveness of infrequent events or to negative valence. This question was addressed in the present study.

In addition to the empirical findings that challenged the original distinctiveness account, alternative theoretical approaches were suggested to explain illusory correlations. One of these models interprets illusory correlation as an effect of information loss and set size in category learning (Fiedler, 1991, 1996; Sanbonmatsu, Shavitt, & Gibson, 1994). Rather than relying on differential encoding or availability of individual statements, this account assumes that illusory correlations are driven by the more accurate extraction of the preponderance of desirable behaviors from the larger sample of Group A than from the smaller sample of Group B. Another theoretical approach rests on social categorization principles and assumes that people use the valence of behaviors to differentiate between the target groups in a search for meaning of the group distinction (Berndsen & Spears, 1997; Haslam, McGarty, & Brown, 1996; McGarty & de la Haye, 1997; McGarty, Haslam, Turner, & Oakes, 1993). In the present study, we investigated whether such general principles of category acquisition can accommodate both illusory and spurious correlations.

### SPURIOUS CORRELATIONS

In a demonstration of spurious correlations in stereotype formation, Schaller and O'Brien (1992) presented a series of statements about anagram solutions by members of two groups labeled 'Group A' and 'Group B'. Each statement contained group membership, the anagram problem, and the outcome of the trial. The anagrams differed in difficulty, as was discernible from the number of letters: Some anagrams consisted of five letters, whereas others consisted of seven letters. For both types of anagrams, Group A showed a higher success rate (i.e., 100% and 25%) than Group B (i.e., 75% and 0%), indicating a higher level of ability for Group A. Reflecting task difficulty, the proportion of successful trials was lower for seven-letter anagrams than for five-letter anagrams in both groups. Moreover, members of Group A were more likely to work on seven-letter anagrams, whereas members of Group B were more likely to work on five-letter anagrams. Although Group A showed a higher success rate than Group B for both easy and difficult anagrams, aggregation across problem type led to a reversed contingency, with a higher overall success rate for Group B (i.e., 60%) than Group A (i.e., 40%). This reversed contingency was spurious in nature, caused by the moderating role of task difficulty for both group membership and success. The stimulus design thus contained an instance of 'Simpson's Paradox' (Simpson, 1951), that is, an incompatibility between the relations of group membership and success at the level of task difficulty and at the aggregate level (i.e., collapsed across task difficulty).

After stimulus presentation, participants judged Group B as superior to Group A in anagram solving ability and general verbal intelligence. In other words, a stereotype emerged that paralleled the spurious correlation between group membership and success rate collapsed across task difficulty. The stereotype was interpreted in terms of an incomplete statistical reasoning process that does not take into account the confounding variable of difficulty (Schaller, 1994; Schaller & O'Brien, 1992). According to this account, the overall covariation between group membership and success was detected, but the moderating role of task difficulty went unnoticed because of a failure to engage in complex reasoning strategies. This explanation was corroborated by studies which demonstrated that explicit instructions to take task difficulty into account (Schaller, 1992b; Schaller & O'Brien, 1992) or training in statistical reasoning with confounding factors (Schaller, Asp, Rosell, & Heim, 1996) prompted impression formation strategies that included task difficulty and led to less biased group judgments. Likewise, motivational concerns to avoid a negative impression of one's own group were shown to increase the complexity of the inference process and to affect the resulting judgments (Schaller, 1992a).

More recent research on stereotype formation in the case of Simpson's Paradox has indicated, however, that participants can extract pairwise contingencies with a confounding factor and nonetheless form erroneous stereotypes reflecting spurious correlations (Fiedler, Walther, Freytag, & Nickel, 2003; Fiedler, Walther, Freytag, & Stryczek, 2002; Meiser, 2003; Meiser & Hewstone, 2004). Hence, other cognitive processes than simplistic reasoning may contribute to the perception of spurious correlations. The aim of our present approach therefore was to analyze the extraction of contingencies with the context factor in more detail and to outline a theoretical account that may jointly accommodate true contingency learning, spurious correlations, and illusory correlations.

### **AN EXEMPLAR-BASED LEARNING FRAMEWORK FOR ILLUSORY AND SPURIOUS CORRELATIONS**

The procedural and empirical similarities between the paradigms of illusory and spurious correlations in stereotype formation call for an attempt to explain the two phenomena by a set of common principles. Moreover, the original explanations of illusory and spurious correlations, which resulted from an isolated investigation of each phenomenon, are limited in their ability to account for the recent findings that challenged the role of distinctiveness and simplistic reasoning. An exemplar-based learning approach that uses common principles of parallel distributed memory models may serve as a theoretical framework that includes illusory and spurious correlations under a joint umbrella and that avoids paradigm-specific explanations. In particular, exemplar-based learning models do not require enhanced memory for individual events, as is assumed by the distinctiveness account of illusory correlations, nor a specific neglect or loss of information concerning the context factor in judgment formation, as is assumed by the account of spurious correlations in terms of simplistic reasoning.

Parallel distributed memory models rest on two main assumptions. First, it is assumed that category information is not represented in terms of a singular memory trace, such as a unitary prototype or schema, but that the characteristics of the category are derived from a set of category exemplars (e.g., Hintzman, 1986; McClelland & Rumelhart, 1985; Rumelhart, Smolensky, McClelland, & Hinton, 1986). Because each exemplar is a noisy realization of the category and thus shows the shared characteristics of the category members with some degree of distortion, aggregation across a reasonable number of exemplars yields the central features of the category as a whole. The second core assumption of parallel distributed memory models is that the memory representation of individual category exemplars is distributed over a number of information components, or units in a connectionist network. As a consequence, the memory traces of individual category exemplars can be conceived of as vectors. The vectorial exemplar representation contains components that denote category membership and components that reflect the features of the exemplar, including those features that are shared by other category members. Given the suppositions of exemplar-based processing and distributed exemplar representation, the total category information in memory can be summarized in a matrix of information components by exemplars (see Fiedler, 1996; Hintzman, 1986). Judgment formation can then be simulated on the basis of the stored exemplar information by probing the memory matrix with a category name. Thus, parallel distributed memory models reflect very general processes of category learning and categorical judgment that are not confined to particular contents or features of the stimulus distribution.

Previous studies have shown that the biased stereotype that is obtained with the usual two-dimensional stimulus distribution of group membership and desirability in the illusory correlation paradigm can be produced by simulations with exemplar-based learning models using parallel distributed memory principles (Fiedler, 1996, 2000; Smith, 1991). Here we investigated whether

such learning models may provide a coherent framework for the integration of illusory correlations, spurious correlations and true contingency learning on the basis of the more complex stimulus distributions used in our present experiment.

## OVERVIEW OF THE PRESENT RESEARCH

To summarize, illusory correlations and spurious correlations are robust phenomena that have been demonstrated in numerous studies on stereotype formation. It has also become clear that their original explanations rely on rather specific characteristics of the stimulus designs used in either paradigm, that is, on a memory effect of the relative infrequency of one group and one class of behaviors in the case of illusory correlations and on the neglect of the confounding role of the context factor in the case of spurious correlations. To advance a more integrative view of cognitive processes in stereotype formation, we investigated illusory correlations and spurious correlations in a joint approach.

For this purpose, we first conducted an experiment to analyze stereotype formation on the basis of different stimulus distributions. The stimulus distributions had the same trivariate structure and were generated by the binary variables of group membership (Group A vs. B), valence of behavior (desirable vs. undesirable), and the context variable of town of residence (Town X vs. Y). The stimulus distributions differed, however, with respect to distinctiveness in terms of relative infrequency and with respect to the confounding role of the context variable. The manipulation of relative infrequency and of the confounding role of the context factor allowed us to explore illusory and spurious correlations within the same experimental setting and to conduct focused tests concerning the processes that may underlie stereotype formation. In particular, the different stimulus distributions served mutually as control conditions to test for memory effects of relative infrequency and to test for the detection or neglect of the moderating role of a confounding context variable.

In a second step, we sought to investigate a more comprehensive and unifying perspective on biased stereotype formation that encompasses both illusory and spurious correlations. In a simulation study, we therefore implemented an exemplar-based learning algorithm that extended previous computer models for the simulation of illusory correlations in the classical two-way stimulus design (Fiedler, 1996, 2000; Smith, 2000). With this implementation, we examined whether exemplar-based category learning provides an integrative framework for biased stereotype formation in trivariate stimulus designs that give rise to illusory and spurious correlations.

## EXPERIMENT

The experiment used the two stimulus distributions in Table 1, which shows the frequencies of desirable and undesirable behaviors among members of the Groups A and B within each of the Towns X and Y. The distributions differ with respect to relative infrequency and with respect to the covariations induced by town of residence. First, Distribution (a) contains twice as many desirable as undesirable behaviors and twice as many statements about Group A as Group B. In Distribution (b), in contrast, there are equal numbers of desirable and undesirable behaviors and equal numbers of statements about Group A and Group B. Hence, Group B and undesirable behaviors are relatively infrequent in Distribution (a), but not in Distribution (b). Second, the variables of the stimulus design are completely orthogonal in Distribution (a), whereas Distribution (b) entails a pattern of covariation that constitutes an instance of Simpson's paradox. In Distribution (b), Group A and desirable

Table 1. Trivariate stimulus distributions with relative infrequency and Simpson's Paradox

Behaviors	Town X		Town Y		Aggregate	
	Group A	Group B	Group A	Group B	Group A	Group B
Distribution (a): Relative infrequency						
Desirable	16	8	8	4	24	12
Undesirable	8	4	4	2	12	6
$\Delta P = 0$		$\Delta P = 0$		$\Delta P = 0$		
Distribution (b): Simpson's Paradox						
Desirable	12	6	3	6	15	12
Undesirable	6	3	6	12	12	15
$\Delta P = 0$		$\Delta P = 0$		$\Delta P = 0.11$		

Note:  $\Delta P = p(\text{desirable behavior} \mid \text{Group A}) - p(\text{desirable behavior} \mid \text{Group B})$ .

behaviors are more likely in Town X than Town Y, and Group B and undesirable behaviors are more likely in Town Y than Town X. Because of the covariations of town with both group membership and desirability, the aggregate across the two towns shows a spurious correlation that associates Group A with a higher degree of desirability than Group B.

The  $\Delta P$  index (Allan, 1980) can be used to elucidate the correlation between group membership and desirability on different levels of the trivariate stimulus distribution.  $\Delta P$  is computed by subtracting the probability of desirable behaviors in Group B from the probability of desirable behaviors in Group A. As a consequence,  $\Delta P = 0$  indicates independence of group membership and desirability, whereas  $\Delta P > 0$  and  $\Delta P < 0$  indicate a correlation with a higher or lower probability of desirable behaviors in Group A than Group B, respectively. According to the orthogonal design of stimulus Distribution (a),  $\Delta P = 0$  holds within the subtables for Town X and Town Y as well as in the aggregate table collapsed across the two towns. In Distribution (b), however, a spurious correlation of  $\Delta P = 0.11$  emerges in the aggregate table, despite  $\Delta P = 0$  within each subtable (see Table 1).

Distribution (a) should give rise to an illusory correlation in favor of Group A because of the relative infrequency of Group B and undesirable behaviors. Distribution (b) should produce a spurious correlation in favor of Group A because of the confounding role of town of residence. Hence, a biased stereotype should be observed with both distributions, although different cognitive processes have traditionally been held responsible for biased judgments in the two stimulus conditions. The joint analysis of illusory and spurious correlations in the present experiment allowed focused tests of the underlying processes by means of comparisons between the stimulus conditions.

In accordance with the distinctiveness account of illusory correlations, a memory advantage for the infrequent class of undesirable behaviors has repeatedly been observed and interpreted as evidence for the distinctiveness of rare events (Hamilton et al., 1985; McConnell et al., 1994). However, relative infrequency was confounded with negative valence in these experiments as well as in other experiments on illusory correlations that revealed better memory for undesirable than desirable behaviors (Klauer & Meiser, 2000; Meiser & Hewstone, 2001). A comparison between the two stimulus conditions of the present experiment, with relative infrequency of undesirable behaviors in Distribution (a) but not in Distribution (b), may reveal whether the memory advantage of undesirable behaviors is due to infrequency or whether it reflects a general negativity effect (Pratto & John, 1991; Skowronski & Carlston, 1989) that is independent of numerical infrequency.

Concerning spurious correlations, the two stimulus conditions allow one to test whether the confounding role of town of residence in Distribution (b) is neglected or whether judgments are sensitive to the covariations of town with desirability and group membership. Sensitivity to the



covariations of the confounding factor would be reflected by a stronger evaluative differentiation between Town X and Town Y on the basis of Distribution (b) than Distribution (a), and by overproportional assignments of statements to the combinations of Group A with Town X and Group B with Town Y on the basis of Distribution (b) relative to Distribution (a).

The present experiment thus extends earlier studies on memory effects of infrequency and on the neglect of confounding factors by using different stimulus distributions that can mutually be taken as control conditions for numerical infrequency and for the confounding role of the context variable, respectively.

## Method

### *Participants*

Participants were 70 students from various departments of a British University who were paid £3 for their participation. They were randomly assigned to one of two stimulus conditions: 35 participants were presented with stimulus Distribution (a), and the remaining 35 participants with Distribution (b).

### *Procedure*

For each participant, 54 target behaviors were randomly drawn from a pool of moderately desirable and moderately undesirable behaviors (see Meiser & Hewstone, 2001, for details). During the presentation phase, the behaviors were displayed as sentences that contained different male first names, membership of Group A or B and residence in Town X or Y according to Table 1. The stimulus sentences were presented in random order on a computer monitor. Each sentence appeared for 8.5 seconds, followed by a 1.5 second pause. The participants were informed that the experiment concerned memory for information about individuals and their behaviors, and they were instructed to read each of the sentences carefully.

After the presentation phase, different dependent measures were assessed. The first task consisted of trait ratings, in which the four combinations of Group A and Group B with Town X and Town Y were rated with respect to ten traits. The trait adjectives were selected on the basis of a scale analysis by Rosenberg, Nelson, and Vivekananthan (1968) and included five positive traits (e.g., 'sociable') and five negative traits (e.g., 'unreliable'). For each combination of group and town, participants had to choose a value on a 10-point rating scale ranging from 0 (i.e., the trait 'does not apply at all') to 9 (i.e., the trait 'applies completely').

The second task was an assignment task, in which the 54 target behaviors were presented in random sequence with 54 new distractor behaviors. The distractor behaviors were drawn from the same pool and included the same numbers of desirable and undesirable items as the target behaviors. For each behavior, participants had to decide in a first step whether the behavior had been displayed during the presentation phase (response 'old') or not (response 'new'). If a behavior was classified as 'old,' participants had to decide in a second step whether the behavior referred to an individual in Town X or Town Y, and in a third step whether it referred to a member of Group A or Group B.

The third task required estimations of the numbers of undesirable behaviors that had been presented about the four combinations of the Groups A and B with the Towns X and Y. For this purpose, the total number of statements about each combination was displayed, and participants had to estimate how many of these statements contained undesirable behaviors.

## Results

### Trait Ratings

Rating scores were computed by reversing the ratings on negative traits and averaging across the ten traits. Table 2 displays the mean rating scores for each combination of group membership and town of residence. A 2 (stimulus condition: Distribution (a) vs. Distribution (b))  $\times$  2 (group: Group A vs. Group B)  $\times$  2 (town: Town X vs. Town Y) mixed-model analysis of variance (ANOVA) with repeated measures on the last two factors revealed significant main effects of stimulus condition,  $F(1, 68) = 4.23$ ,  $p = 0.044$ , group,  $F(1, 68) = 7.17$ ,  $p = 0.009$ , and town,  $F(1, 68) = 17.79$ ,  $p < 0.001$ . As can be seen in Table 2, Group A received more positive ratings than Group B, which reflects the expected stereotype in favor of Group A. The effect of group was not moderated by an interaction with stimulus condition,  $F < 1$ , so that the strength of the resulting stereotype did not differ between the stimulus Distributions (a) and (b). The interaction between town and stimulus condition, however, approached significance,  $F(1, 68) = 3.48$ ,  $p = 0.067$ . The interaction between group and town and the higher order interaction were not significant, both  $F < 1$ .

To explore the interaction between town and stimulus condition further, simple effects analyses were conducted for each condition. The analyses revealed that Town X was rated much more positively than Town Y on the basis of Distribution (b),  $F(1, 34) = 15.47$ ,  $p < 0.001$ , whereas the difference between the two towns was only marginally significant for Distribution (a),  $F(1, 34) = 3.44$ ,  $p = 0.072$ . While the marginal effect of town within stimulus condition (a) may reflect an illusory correlation between town and desirability, the significantly stronger evaluative differentiation between Town X and Town Y after presentation of Distribution (b) than Distribution (a) indicates that the actual contingency between town and desirability was extracted from stimulus Distribution (b).<sup>1</sup>

Table 2. Mean trait rating scores and estimated proportions of undesirable behaviors

Stimulus condition	Town X				Town Y			
	Group A		Group B		Group A		Group B	
	<i>M</i>	<i>SD</i>	<i>M</i>	<i>SD</i>	<i>M</i>	<i>SD</i>	<i>M</i>	<i>SD</i>
	Trait rating scores							
Distribution (a)	5.65	1.10	4.67	1.18	5.09	1.20	4.57	1.19
Distribution (b)	5.39	1.44	5.03	1.25	4.55	1.05	4.17	1.56
	Estimated proportions of undesirable behaviors							
Distribution (a)	0.37	0.19	0.57	0.21	0.40	0.21	0.55	0.24
Distribution (b)	0.39	0.20	0.42	0.20	0.53	0.24	0.58	0.19

*Note:* Higher trait rating scores indicate more positive evaluations. Higher proportions of undesirable behaviors indicate more negative evaluations.

<sup>1</sup>The interaction between town and stimulus condition also provides an interpretation for the main effect of stimulus condition. As can be seen in Table 1, Town X is characterized by 67% desirable behaviors in both stimulus distributions, whereas Town Y is characterized by 67% desirable behaviors in Distribution (a) but only 33% desirable behaviors in Distribution (b). Accordingly, it may be expected that judgments of Town X are very similar for both stimulus conditions and that judgments of Town Y differ across conditions. This expectation was confirmed by simple effects analyses of the ratings for each of the towns. While ratings of Town X did not differ between stimulus conditions,  $F < 1$ , ratings for Town Y were more positive on the basis of Distribution (a) than Distribution (b),  $F(1, 68) = 8.08$ ,  $p = 0.006$ . The effect of stimulus condition was thus limited to the ratings of Town Y.



### *Estimated Proportions of Undesirable Behaviors*

Estimated proportions were computed by dividing the estimated numbers of undesirable behaviors by the total number of statements presented about each combination of group membership and town of residence. The mean estimated proportions are reported in Table 2. The ANOVA yielded significant main effects of group,  $F(1, 68) = 17.64$ ,  $p < 0.001$ , and town,  $F(1, 68) = 9.37$ ,  $p = 0.003$ . The effects indicate that lower proportions of undesirable behaviors were estimated for Group A than Group B and for Town X than Town Y. The main effect of stimulus condition did not attain significance,  $F < 1$ , but condition was involved in interactions with group,  $F(1, 68) = 6.61$ ,  $p = 0.012$ , and town,  $F(1, 68) = 8.29$ ,  $p = 0.005$ . There was no interaction between group and town, nor a higher order interaction, both  $F < 1$ .

Simple effects analyses of the interaction between group and stimulus condition revealed that the estimated proportion of undesirable behaviors was lower for Group A than Group B on the basis of stimulus Distribution (a),  $F(1, 34) = 32.28$ ,  $p < 0.001$ , but not on the basis of Distribution (b),  $F(1, 34) = 1.03$ ,  $p = 0.317$ . Thus, the expected stereotype in favor of Group A was evident for Distribution (a), but it failed to reach significance for Distribution (b). While the significant effect of target group within stimulus condition (a) reflects the illusory correlation effect, the reasons for the failure to obtain a significant spurious correlation effect in condition (b) of the present experiment are not clear. Simple effects analyses of the interaction between town and stimulus condition yielded significantly lower estimates of undesirable behaviors for Town X than Town Y on the basis of Distribution (b),  $F(1, 34) = 19.66$ ,  $p < 0.001$ , but not on the basis of Distribution (a),  $F < 1$ . The interaction thereby indicates a stronger evaluative differentiation between Town X and Town Y after presentation of Distribution (b) than Distribution (a). This result suggests that the actual covariation between town and desirability in Distribution (b) was perceived and used to differentiate between the two towns.

### *Assignment Task*

Assignment frequencies were computed from the responses to desirable and undesirable target and distractor behaviors in the assignment task of the two stimulus conditions. There were target items from the four combinations of Town X and Town Y with Group A and Group B, and there were new distractor items. Likewise, there were five response alternatives for each test item, namely assignment to one of the four combinations of town and group and the response 'new.' The cell frequencies in the resulting  $5$  (item type)  $\times 5$  (response category) tables for desirable and undesirable behaviors were analyzed with a multinomial source monitoring model that provides measures of recognition memory, source memory for two dimensions of context information, and guessing rates on various stages of the assignment task (Meiser, 2005; Meiser & Bröder, 2002).

The multinomial source monitoring model is illustrated in Figure 1. Each branch in the figure represents a combination of cognitive states which are evoked by a given test item and which jointly lead to an observable response. The various cognitive states are specified by model parameters. The parameter  $D$  denotes recognition memory for the behaviors, the parameters  $d^{\text{town}}$  and  $d^{\text{group}}$  denote source memory for the town and group context of the behaviors,  $b$  reflects the guessing tendency to classify items as old, and  $g^{\text{town}}$ ,  $g_{\text{X}}^{\text{group}}$  and  $g_{\text{Y}}^{\text{group}}$  specify guessing rates in the assignment of behaviors to the towns and groups. Table 3 provides the exact definitions of the model parameters.

The measurement model in Figure 1 allows an in-depth analysis of episodic memory and reconstructive guessing processes in stereotype formation. For this purpose, separate sets of model parameters were specified for the responses to desirable and undesirable behaviors in each stimulus




Figure 1. Processing tree diagram of the multinomial source monitoring model. Town  $i \in \{X, Y\}$ , Group  $j \in \{A, B\}$ . The model parameters are described in Table 3. Adapted from 'Memory for Multidimensional Source Information' by T. Meiser and A. Bröder, 2002, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 28, pp. 116–137

condition. To test the overall goodness of model fit, we used the likelihood ratio statistic  $G^2$  which asymptotically approaches a  $\chi^2$  distribution (Batchelder & Riefer, 1999). The model showed an excellent fit to the empirical frequency data,  $G^2(52) = 44.29$ ,  $p = 0.768$ , and could thus be maintained as a valid model of the cognitive processes in the assignment task. Table 4 shows the estimates of the

Table 3. Model parameters in the multinomial source monitoring model

Parameter	Definition
$D$	Probability of recognizing target behaviors as old and distractor behaviors as new
$d^{\text{town}}$	Probability of remembering the town origin of a recognized target behavior
$d^{\text{group}}$	Probability of remembering the group origin of a recognized target behavior
$b$	Probability of guessing that a nonrecognized behavior is old
$g^{\text{town}}$	Probability of guessing that a behavior referred to Town X
$g_{ X}^{\text{group}}$	Probability of guessing that a behavior referred to Group A given assignment to Town X
$g_{ Y}^{\text{group}}$	Probability of guessing that a behavior referred to Group A given assignment to Town Y

model parameters. To compare the model parameters between the stimulus conditions and between the assignment of desirable and undesirable behaviors, we used the conditional likelihood ratio statistic  $\Delta G^2$ , which follows a  $\chi^2$  distribution with one degree of freedom (Batchelder & Riefer, 1999). The right-hand column of Table 4 shows the results of pairwise parameter comparisons between the stimulus conditions. Parameter estimates that differ significantly between responses to desirable versus undesirable behaviors are marked by an asterisk.

Concerning memory performance for the different kinds of stimuli, recognition memory  $D$  was significantly better for undesirable than desirable behaviors. This memory advantage for undesirable behaviors was found for both stimulus distributions, and there were no differences in recognition memory between stimulus conditions either for desirable or undesirable items. The data thus revealed a memory advantage for undesirable behaviors that was equal for Distribution (a), in which undesirable behaviors were infrequent, and Distribution (b), in which no kind of information was

Table 4. Parameter estimates, 95% confidence intervals (CI), and tests for equality of parameters in the source monitoring model

Parameter	Distribution (a)		Distribution (b)		$\Delta G^2(1)$	$p$
	Estimate	CI	Estimate	CI		
Desirable behaviors						
$D$	0.35*	(0.31, 0.38)	0.38*	(0.34, 0.42)	1.11	0.292
$b$	0.34*	(0.31, 0.37)	0.32	(0.29, 0.35)	0.53	0.468
$d^{\text{town}}$	0.00	(−0.12, 0.12)	0.06	(−0.07, 0.19)	0.59	0.442
$d^{\text{group}}$	0.10	(−0.03, 0.22)	0.00	(−0.13, 0.13)	1.35	0.245
$g^{\text{town}}$	0.48	(0.44, 0.51)	0.59*	(0.55, 0.63)	19.89	< 0.001
$g_{ X}^{\text{group}}$	0.55*	(0.50, 0.60)	0.62*	(0.57, 0.66)	3.69	0.055
$g_{ Y}^{\text{group}}$	0.59*	(0.54, 0.63)	0.52*	(0.47, 0.58)	2.68	0.101
Undesirable behaviors						
$D$	0.55*	(0.51, 0.60)	0.55*	(0.52, 0.59)	0.00	0.986
$b$	0.27*	(0.23, 0.32)	0.32	(0.28, 0.36)	2.37	0.123
$d^{\text{town}}$	0.00	(−0.12, 0.12)	0.00	(−0.10, 0.10)	0.00	0.999
$d^{\text{group}}$	0.01	(−0.11, 0.13)	0.08	(−0.01, 0.18)	0.83	0.361
$g^{\text{town}}$	0.50	(0.45, 0.55)	0.41*	(0.38, 0.45)	8.31	0.004
$g_{ X}^{\text{group}}$	0.38*	(0.32, 0.45)	0.45*	(0.40, 0.51)	2.63	0.105
$g_{ Y}^{\text{group}}$	0.49*	(0.43, 0.56)	0.41*	(0.37, 0.46)	3.59	0.058

*Note:* The values of the conditional likelihood ratio statistic  $\Delta G^2(1)$  refer to tests for equality of model parameters between stimulus conditions. Model parameters that differ significantly between responses to desirable and undesirable behaviors are marked with an asterisk, all  $\Delta G^2(1) > 3.84$ ,  $p < 0.05$ .

infrequent. The observed memory advantage for undesirable behaviors therefore cannot be caused by the distinctiveness of infrequent events. Instead, it appears to reflect a general negativity effect, such as enhanced processing and retention of negative information independent of infrequency (Pratto & John, 1991; Skowronski & Carlston, 1989).<sup>2</sup>

Unlike recognition performance, source memory for the town and group origin of target behaviors,  $d^{\text{town}}$  and  $d^{\text{group}}$ , was consistently low. In line with previous results (Meiser & Hewstone, 2001), all source memory parameters were close to zero (see parameter estimates and confidence intervals in Table 4). There were no differences in source memory performance for desirable and undesirable behaviors or between stimulus conditions.<sup>3</sup>

Turning to the guessing parameters, earlier research showed that guessing processes in source attributions reflect perceived contingencies between behavioral characteristics and social categories (Bayen, Nakamura, Dupuis, & Yang, 2000; Klauer & Meiser, 2000; Meiser, 2003; Meiser & Hewstone, 2004; Wegener & Klauer, 2004). In line with this reasoning, the guessing parameter for the attribution of behaviors to Town X versus Y,  $g^{\text{town}}$ , showed a pattern that resembled the true existence or nonexistence of a contingency between town and desirability. For Distribution (a),  $g^{\text{town}}$  did not differ between desirable and undesirable behaviors. For Distribution (b), in contrast,  $g^{\text{town}}$  was significantly larger for desirable than undesirable behaviors, indicating that desirable behaviors were more likely to be attributed to Town X than were undesirable behaviors. Moreover, comparisons between stimulus conditions revealed that  $g^{\text{town}}$  was smaller for Distribution (a) than Distribution (b) in the case of desirable behaviors and that the opposite difference emerged in the case of undesirable behaviors. Together, the effects indicate that desirability was effectively used as a cue for town assignments on the basis of Distribution (b) but not Distribution (a), corresponding to the true relations between town and desirability in the stimuli. These findings match the results of the trait ratings and frequency estimates in revealing that the confounding role of town of residence in Distribution (b) was perceived and taken into account.

The guessing parameters for the attribution of behaviors to Group A versus B,  $g_{|X}^{\text{group}}$  and  $g_{|Y}^{\text{group}}$ , indicated the expected group stereotype in favor of Group A. The parameters  $g_{|X}^{\text{group}}$  and  $g_{|Y}^{\text{group}}$  were significantly larger for desirable than undesirable behaviors in both stimulus conditions. These results show that desirable behaviors were consistently more likely to be attributed to Group A than were undesirable behaviors, irrespective of their prior assignment to Town X or Town Y and irrespective of

<sup>2</sup>Additional tests of recognition performance  $D$  were carried out to analyze whether recognition memory varied as a function of the town or group origin of target behaviors. The tests showed no differences in  $D$  between the four combinations of town and group for desirable or undesirable items in the two stimulus conditions, all  $\Delta G^2(3) < 4.45, p > 0.217$ . These results also imply that the relative infrequency of Group B and Town Y in Distribution (a) had no effects on recognition performance. The findings of the multinomial model were confirmed by a conventional signal detection analysis of recognition performance (see Snodgrass & Corwin, 1988). The index of memory accuracy  $d'$  was derived from the hit and false alarm rates in the assignment task. A  $2$  (stimulus condition)  $\times 2$  (town)  $\times 2$  (group)  $\times 2$  (desirability) ANOVA with repeated measures on the last three factors revealed significantly higher values of  $d'$  for undesirable than desirable behaviors,  $F(1, 68) = 78.19, p < 0.001$ . Importantly, this memory advantage for undesirable behaviors was not qualified by an interaction with stimulus condition,  $F < 1$ , indicating that memory accuracy was not affected by the relative infrequency of undesirable behaviors in Distribution (a). The only further effect that attained significance was the three-factor interaction of desirability with town and stimulus condition,  $F(1, 68) = 5.98, p = 0.017$ . The interaction is due to the fact that the memory advantage for undesirable items was somewhat stronger for behaviors from Town X in Distribution (a), whereas it was somewhat stronger for behaviors from Town Y in Distribution (b). However, simple effects analyses ensured that there was no difference in  $d'$  between stimulus conditions for undesirable behaviors from Town X,  $F(1, 68) = 1.17, p = 0.283$ , nor for undesirable behaviors from Town Y,  $F < 1$ . Likewise, there was no difference in  $d'$  between stimulus conditions for desirable behaviors from Town X or Town Y, both  $F < 1$ .

<sup>3</sup>Additional tests were performed to analyze whether source memory for town and group,  $d^{\text{town}}$  and  $d^{\text{group}}$ , varied as a function of the actual town and group origin of target behaviors. The tests revealed no differences in source memory for town, all  $\Delta G^2(3) < 6.35, p > 0.096$ , or in source memory for group, all  $\Delta G^2(3) < 2.51, p > .474$ , for desirable and undesirable items in both stimulus conditions.

stimulus condition.<sup>4</sup> This overproportional assignment of desirable behaviors to Group A relative to undesirable behaviors, which contradicts the actual independence of group membership and desirability in each town context of the stimulus distribution (see Table 1), mirrors the erroneous preference for Group A that was also observed in the trait ratings and frequency estimates.

Finally, the guessing parameters  $g_{|X}^{\text{group}}$  and  $g_{|Y}^{\text{group}}$  allowed us to test for a perceived contingency between town of residence and group membership. In particular, we expected an overproportional assignment of behaviors to the combinations of Group A with Town X and Group B with Town Y after presentation of stimulus Distribution (b), indicating sensitivity to the actual covariation of the towns and groups. The comparisons of the parameter  $g_{|X}^{\text{group}}$  between the two stimulus conditions showed that behaviors assigned to Town X were more likely to be attributed to Group A on the basis of Distribution (b) than Distribution (a) (see Table 4). Analogously, the comparisons of  $g_{|Y}^{\text{group}}$  showed that behaviors assigned to Town Y were less likely to be attributed to Group A (i.e., more likely to be attributed to Group B) on the basis of Distribution (b) than Distribution (a). Although these parameter comparisons reached only marginal levels of significance, the parameter estimates suggest stronger associations of Town X with Group A and Town Y with Group B after presentation of Distribution (b), reflecting the actual degree of correlation in the stimuli.

## Discussion

To summarize, a stereotype in favor of Group A was revealed by the trait ratings, the estimated proportions of undesirable behaviors, and the guessing processes in behavior assignments. The stereotype corresponds to an illusory correlation for stimulus Distribution (a), with relative infrequency of Group B and undesirable behaviors, and to a spurious correlation for Distribution (b), with a misleading contingency between group membership and desirability in the aggregate table.

Despite the evidence of an illusory correlation on the basis of Distribution (a), there was no indication of a memory effect of relative infrequency, which replicates and extends previous results (Fiedler et al., 1993; Klauer & Meiser, 2000; Meiser & Hewstone, 2001). In particular, the present experiment yielded original evidence that the memory advantage for undesirable behaviors, which had been interpreted as support for the distinctiveness account of illusory correlations (e.g., Hamilton et al., 1985; McConnell et al., 1994), can be traced back to a general negativity effect independent of numerical infrequency. Because the memory advantage for undesirable behaviors was not influenced by the infrequency manipulation, and because the infrequency of Group B in Distribution (a) did not affect any of the memory measures, the present results add further evidence that enhanced memory for individual items, such as negative and/or relatively infrequent events, is neither specific nor causal to the formation of illusory correlations.

Concerning the spurious correlation on the basis of Distribution (b), the comparisons within and across stimulus conditions showed that participants were quite sensitive to the covariations with town

<sup>4</sup>The results and interpretations of the guessing parameters  $g^{\text{town}}$ ,  $g_{|X}^{\text{group}}$  and  $g_{|Y}^{\text{group}}$  were corroborated in a traditional loglinear analysis of source attributions for distractor items that were mistakenly judged "old" (i.e., false alarms). A hierarchical loglinear analysis of the 2 (town)  $\times$  2 (group)  $\times$  2 (desirability)  $\times$  2 (stimulus condition) table of false alarms exhibited a significant three factor interaction of town attributions and desirability with stimulus condition,  $G^2(1) = 5.26, p = 0.022$ . In line with the guessing parameter  $g^{\text{town}}$  in Table 4, the loglinear three factor interaction showed that town attributions of false alarms were stochastically dependent on desirability for Distribution (b),  $G^2(1) = 13.97, p < 0.001$ , whereas town attributions were independent of desirability for Distribution (a),  $G^2(1) = 0.05, p = 0.821$ . The loglinear analysis further showed that there was no three factor interaction of group attributions and desirability with stimulus condition,  $G^2(1) = 0.39, p = 0.534$ . The absence of this interaction confirms that the overproportional assignment of desirable behaviors to Group A was identical for the two stimulus distributions, which supports the conclusions that were drawn from the results of the guessing parameters  $g_{|X}^{\text{group}}$  and  $g_{|Y}^{\text{group}}$ .

of residence, rather than ignoring the context factor. The pronounced evaluative differentiation between Town X and Town Y that was obtained in all dependent measures after presentation of Distribution (b) reflected the actual covariation between town and desirability in the stimuli. Moreover, the guessing processes in group assignments indicated that the covariation between the towns and groups was also discerned. Thereby, the results of the guessing parameters revealed that participants did not give independent judgments about the two groups on the one hand and the two towns on the other, but that they simultaneously used the covariations of the context factor with both group membership and desirability for their responses.

The present findings question the necessity of rather specific assumptions about stereotype formation in the case of relative infrequency and in the case of a confounding context factor, respectively. Instead, the results encourage one to search for a common theoretical framework for illusory and spurious correlations that does not rest on enhanced memory for rare events nor on the neglect of a confounding context factor. As mentioned in the Introduction, general principles of information loss (Fiedler, 1991, 1996; Sanbonmatsu et al., 1994) and social categorization (Berndsen & Spears, 1997; McGarty & de la Haye, 1997) may provide overarching perspectives that parsimoniously accommodate biased stereotype formation in both paradigms, and such general principles can be captured in exemplar-based models of category learning. In the following study, we therefore derived the implications of an exemplar-based learning algorithm concerning illusory and spurious correlation by means of computer simulation.

## SIMULATION STUDY

The results of the experiment suggest some features that a model should have if it is to explain erroneous stereotype formation in the present experimental setting. A successful model should account not only for biased stereotypes but also for the extraction of existing covariations with the confounding factor. Given the extremely poor source memory for the town and group origin of all types of behaviors and the lack of any infrequency-based memory effect, the model should not rely on enhanced episodic memory for individual events that link a certain group to a certain kind of behavior. Finally, an ideal model should account for erroneous stereotype formation on the basis of both stimulus Distributions (a) and (b) in Table 1. By accommodating stereotypes that correspond to illusory correlations and spurious correlations alike, such a model can help to overcome paradigm-specific explanations that have largely dominated the literature so far. Here we will demonstrate that an exemplar-based learning algorithm that adopts parallel distributed memory principles of category acquisition can meet these criteria.

### Implementation of a Modified BIAS Model

To simulate stereotype formation, we implemented a modified version of the BIAS (Brunswikian Induction Algorithm for Social Cognition; Fiedler, 1996) algorithm.<sup>5</sup> Figure 2 illustrates the matrix representation of the 54 exemplars in the stimulus distributions of Table 1. Each column of the matrix represents the memory trace for an individual statement referring to a member of Group A or B who is

<sup>5</sup>BIAS is an exemplar-based learning model that simulates impression formation by information aggregation. In contrast to connectionist network models, BIAS simulates learning by adding vectors to the information matrix, rather than recalculating connection strengths between processing nodes.



		Exemplars						
		Group A Town X desirable	Group A Town Y desirable	Group B Town X desirable	Group A Town X undesirable	...	Group B Town Y undesirable	Group A Town X desirable
Group segment	{	+1	+1	-1	+1	...	-1	+1
		-1	-1	+1	-1		+1	-1
		-1	-1	+1	-1		+1	-1
		+1	+1	-1	+1	...	-1	+1
Town segment	{	-1	+1	-1	-1	...	+1	-1
		-1	+1	-1	-1		+1	-1
		+1	-1	+1	+1		-1	+1
		+1	-1	+1	+1	...	-1	+1
Desirability segment	{	+1	+1	+1	-1	...	-1	+1
		-1	-1	-1	+1		+1	-1
		+1	+1	+1	-1		-1	+1
		-1	-1	-1	+1		+1	-1
		-1	-1	-1	+1		+1	-1
		+1	+1	+1	-1		-1	+1
		+1	+1	+1	-1		-1	+1
		-1	-1	-1	+1	...	+1	-1
Specifics (e.g., actor, behavior, context)	{	+1	+1	-1	+1	...	-1	-1
		-1	+1	-1	+1		+1	-1
		+1	-1	-1	-1		+1	+1
		+1	-1	+1	-1		+1	-1
		+1	+1	-1	-1		-1	-1
		⋮	⋮	⋮	⋮		⋮	⋮
		-1	-1	-1	+1		+1	-1
		+1	-1	+1	+1	...	-1	+1

Figure 2. Matrix representation of the stimulus exemplars in the simulation model

living in Town X or Y and behaves in a desirable or undesirable way. For convenience, each memory trace consists of 30 components. Eight components denote the combination of group membership and town of residence, another eight components indicate the desirability of the behavior, and the remaining components represent specifics of the individual statement, such as the actor's name, particulars of the behavior, or any information concerning the situational context of stimulus presentation.

The matrix in Figure 2 contains the true vectors of the 54 stimuli, that is, stimulus information without distortion. For the following simulations, varying degrees of distortion were specified in the memory representations of the simulated 'participants.' The distortion process captures various kinds of noise and information loss, such as interpretive uncertainty regarding the stimuli, attention deficits during encoding, memory decay, or retrieval difficulties at the time of judgment. Three distortion

parameters were implemented in the simulation model. The first distortion parameter concerned the vector components of group membership, the second parameter concerned the components of town of residence, and the third parameter concerned the components referring to the behavior information, including desirability. The specification of different distortion processes was motivated by prevailing evidence that memory for events can be dissociated from source memory for context information (Bayen, Murnane, & Erdfelder, 1996; Jurica & Shimamura, 1999) and that source memory for different context dimensions can be dissociated from each other (Meiser & Bröder, 2002; Meiser & Hewstone, 2004).

The distortion parameters introduced a probabilistic and componentwise error process in the stimulus matrix. Each component of a given vector segment was independently replaced by a random value of  $-1$  or  $+1$  with a certain probability. The three distortion parameters were varied orthogonally in three steps. Low distortion was defined as a 25% probability that a component is replaced by a random value. Medium distortion was defined as a 50% probability, and high distortion as a 75% probability of replacement by a random value. Because the distortion processes affected all stimulus vectors of the memory matrix in the same way, the simulation model did not specify enhanced memory (i.e., reduced distortion) for any class of stimuli, such as infrequent or paired-infrequent events. By the orthogonal variation of the degrees of distortion concerning group membership, town of residence, and desirability, the model included information loss for all aspects of the trivariate stimuli, rather than selective forgetting or neglect of the confounding context factor. The simulation model therefore contained general processes of information loss and avoided specific assumptions akin to distinctiveness or simplistic reasoning.


Impression formation was simulated by prompting the distorted stimulus matrix with the four composite vector segments that represented the combinations of a particular group with a particular town. The prompting process results in an aggregate vector for a given category that reflects the central tendency of the stored category exemplars. First, each column vector was weighted by its dot product with the prompt vector. Since all matrix elements were either  $+1$  or  $-1$ , the dot product is equal to the number of matches minus the number of mismatches between the prompt vector and the distorted group and town segments of the memory trace. Second, if the weight was positive, the weighted memory trace was added to the aggregate vector of the particular combination of group and town. Third, following aggregation across memory traces, correlations were computed between the desirability segments of the four aggregate vectors and the ideal segment of desirable behaviors. The correlations indicate the degree to which a given combination of group membership and town of residence is associated with desirability in the case of a positive correlation, and with undesirability in the case of a negative correlation.

### **Simulated Group Stereotypes**

Simulations were run with 500 'participants' for each constellation of distortion parameters. Figure 3 displays the resulting correlations with desirability for the four combinations of group membership and town of residence on the basis of the stimulus Distributions (a) and (b).

As can be seen in Figure 3, Group A in Town X (black bars) shows higher correlations with desirability than Group B in Town X (dark grey bars) across the two stimulus distributions and across a wide range of distortion parameters. For Distribution (a), Group A in Town Y (light grey bars) also shows stronger correlations with desirability than Group B in Town Y (white bars). For Distribution (b), correlations involving Town Y are negative, which indicates associations with undesirability and reflects the actual majority of undesirable behaviors in Town Y (see Table 1). Nonetheless, Group A in Town Y shows less negative correlations with desirability than does Group B in Town Y. The

## Stimulus Distribution (a)



## Stimulus Distribution (b)




Figure 3. Mean correlations with desirability in the computer simulations of stereotype formation for the stimulus Distributions (a) and (b). Black bars represent Group A in Town X, dark grey bars represent Group B in Town X, light grey bars represent Group A in Town Y, and white bars represent Group B in Town Y. 'l' indicates low distortion (i.e., each component of a segment is replaced by a random value with a probability of 25%), 'm' indicates medium distortion (i.e., probability of 50%), and 'h' indicates high distortion (i.e., probability of 75%)

simulated judgments thereby reflect a group stereotype in favor of Group A in terms of stronger associations with desirability and weaker associations with undesirability for Group A than Group B. At the same time, the simulated judgments mirror the existing covariation between town and desirability in stimulus Distribution (b) in terms of positive correlations for Town X as opposed to negative correlations for Town Y.

Taken together, the BIAS algorithm produced stereotypes in favor of Group A corresponding to illusory and spurious correlations, and it extracted the actual contingency between town of residence and desirability in stimulus Distribution (b). The simulation results resemble the group stereotype and the differential evaluations of the two towns that were observed in the previous experiment. Importantly, biased stereotype formation was simulated on the basis of general principles of memory storage, information loss, and aggregation over distorted category exemplars, without assuming processes that were specifically tailored to the characteristics of either of the stimulus distributions. Instead, the simulated group stereotypes merely reflect better extraction of the proportions of desirable and undesirable behaviors from larger samples of distorted exemplars as compared with smaller samples. Concerning Distribution (a), the predominance of desirable behaviors is reproduced more accurately for Group A than Group B, because the sample of Group A members is larger than the sample of Group B members within both towns. Concerning Distribution (b), the majority of desirable behaviors in Town X is better extracted from the larger combination of Town X with Group A, whereas the majority of undesirable behaviors in Town Y is better extracted from the larger combination of Town Y with Group B. Thus, according to the exemplar-based category learning account, the biased group stereotypes result as a natural consequence of aggregation over samples of

different size, regardless of whether the stimulus distribution contains paired infrequent stimuli or Simpson's Paradox.

## GENERAL DISCUSSION

The present research pursued an integrative perspective on the cognitive processes underlying erroneous stereotype formation in different experimental paradigms. Hitherto, the paradigms of distinctiveness-based illusory correlations (Hamilton & Gifford, 1976) and spurious correlations (Schaller & O'Brien, 1992) have been treated as distinct entities with respect to their empirical findings and the theoretical explanations of biased stereotype formation. In a first step, we therefore analyzed illusory and spurious correlations in a joint experimental setting by manipulating relative infrequency and the moderating role of a context variable in a trivariate stimulus design. Although the expected stereotype was found, there was no effect of relative infrequency on memory for the behavioral statements or their social origin, nor did participants ignore existing covariations with a confounding context factor. Instead, comparisons across the stimulus conditions revealed that enhanced memory for undesirable behaviors was due to valence, rather than infrequency, and that participants were quite accurate in extracting the moderating role of the context factor. These findings question the necessity of cognitive processing assumptions that focus on specific features of the stimulus distribution to explain biased stereotypes.

Concerning the original distinctiveness account of illusory correlations, the assumption of enhanced memory for information linking members of the minority group to infrequent undesirable behaviors was questioned by the observations that none of the memory parameters in the multinomial model varied as a function of relative infrequency and that source memory for group membership was consistently close to chance level (see also Klauer & Meiser, 2000; Meiser & Hewstone, 2001). Especially, the present experiment showed that better recognition memory for undesirable than desirable behaviors remains unchanged when undesirability is unraveled from relative infrequency. The memory advantage of undesirable behaviors, which had previously been interpreted as evidence for the distinctiveness of infrequent behaviors in the illusory correlation paradigm (Hamilton et al., 1985; McConnell et al., 1994), therefore appears to reflect an effect of valence rather than infrequency. Notwithstanding the lack of an infrequency effect on memory accuracy, however, the present results do not rule out the possibility that distinctiveness operates via subjective memory experience, such as perceived fluency of paired infrequent behaviors during retrieval (e.g., Stroessner & Plaks, 2001). Future research should therefore go beyond measures of objective memory performance and include assessments of subjective retrieval experience to investigate the role of distinctiveness.

Concerning the original account of spurious correlations in terms of incomplete statistical reasoning, a biased stereotype was found in the present experiment, although participants proved to be clearly sensitive to the confounding role of town of residence in stimulus Distribution (b). As shown by the comparison between stimulus conditions, the evaluative differentiation between the two towns reflected responsiveness to the actual covariation between town of residence and desirability, and the behavior assignments also indicated awareness of the covariation between town of residence and group membership. Thus, the stereotype that was observed in the stimulus condition with Simpson's Paradox seems not to result from an incomplete reasoning strategy of overlooking the role of the context variable. Because the present experimental procedures differed from the procedures used in earlier research on spurious correlations (e.g., Schaller, 1992a, 1992b; Schaller & O'Brien, 1992) in various ways (see Meiser & Hewstone, 2004, for a detailed discussion), our conclusions are not intended to criticize the tenability of the simplistic reasoning account for the earlier findings. Instead,

our results imply that neglect of a confounding context variable is not a necessary prerequisite for spurious correlations to be formed, although such neglect may cause spurious correlations under different experimental conditions.

Given that the present findings did not match theoretical accounts that focus on specific characteristics of the stimulus distributions used in the paradigms of illusory and spurious correlations, respectively, we proposed an alternative account that accommodates stereotype formation for different stimulus distributions on the basis of quite general processing assumptions. For this purpose, we implemented an extension of the BIAS learning algorithm (Fiedler, 1996). The model specified an exemplar-based and distributed representation of social information that is subject to information loss of varying degrees. Judgments were modeled by aggregation over the stored exemplar information for each category and by comparing the resulting aggregates with an ideal vector of desirability. Computer simulations demonstrated that the model produces biased stereotypes corresponding to illusory and spurious correlations together with true contingency learning with respect to the confounding context factor. The simulation results thereby closely matched the empirical findings in our experiment on illusory and spurious correlations. Moreover, the present simulation of a spurious correlation on the basis of stimulus Distribution (b) generalizes a previous demonstration that a modified BIAS algorithm can reproduce empirically observed biases in a different stimulus design with a confounding context factor (Meiser & Hewstone, 2004). Although the BIAS model implemented here may not account for all specific results obtained in the spurious correlation paradigm (see Meiser & Hewstone, 2004, Study 2), it accommodates the general outcomes of illusory and spurious correlation in an integrated theoretical framework.

In the exemplar-based learning model, biased stereotypes occur as natural effects of aggregating over different set sizes in a probabilistic, error-prone environment. Different set sizes can result from skewed overall frequencies, as in the illusory correlation paradigm, from covariations that produce more frequent and less frequent category combinations, as in the spurious correlation paradigm, from biased sampling of social information, and so on. Moreover, given that most social information is not perfectly reliable and that the capabilities of human learning and memory are limited, information loss is an ubiquitous feature of social information processing. According to exemplar-based learning models, the co-existence of different set sizes and information loss is sufficient to produce judgment biases inasmuch as the true characteristics of a category are extracted more accurately from a larger sample of noisy exemplars than from a smaller sample of noisy exemplars.

In the present modeling approach, this implication was confirmed for trivariate stimulus designs containing an illusory correlation or a spurious correlation (see Figure 3). With respect to Distribution (a) that produces an illusory correlation, the simulations showed better extraction of the predominance of desirable behaviors for the larger categories involving Group A than for the smaller categories involving Group B (see also Fiedler, 1996, 2000; Smith, 1991). With respect to Distribution (b) that produces a spurious correlation, the simulations revealed better extraction of the true predominance of either desirable or undesirable behaviors for the larger categories involving either Group A (i.e., in Town X) or Group B (i.e., in Town Y). Thus, no matter whether different set sizes result from overall infrequency or from covariations in the stimuli, the fact that categories differ in size, together with the prevalence of information loss, proved sufficient to engender judgment biases which resembled the illusory correlations and spurious correlations that are typically observed in human experiments.

The simulation results highlight that illusory correlations and spurious correlations may be accounted for by common processes of exemplar-based category learning. In line with theoretical approaches that are based on broader principles of social categorization (e.g., Berndsen & Spears, 1997; McGarty & de la Haye, 1997), behavioral valence is used to differentiate between the artificial groups in the experimental setting, and the valence information is differentially extracted from the categories because of their different size. The computer simulations thereby showed that the

exemplar-based learning account integrates the as yet separate phenomena of illusory and spurious correlations. The integration of empirical phenomena into a coherent and parsimonious framework is one of the major theoretical advantages of computer simulations in social-psychological theorizing (Smith, 1996) that allows one to analyze basic cognitive processes which lead to seemingly different effects in various experimental paradigms (Smith & DeCoster, 1998; van Rooy, van Overwalle, Vanhooissen, Labiouse, & French, 2003). In line with this notion, the empirical results of the present experiment and the simulation results with the exemplar-based learning model suggest a joint theoretical perspective on erroneous group stereotype formation and contingency learning in different paradigms in terms of a common set of elementary cognitive processes that were implemented in the exemplar-based learning algorithm.

### ACKNOWLEDGMENTS

The experiment reported in this article was conducted while the first author was a Feodor Lynen Research Fellow at Cardiff University/UK, sponsored by the Alexander von Humboldt Foundation. The computer simulations and manuscript preparation were supported by a grant from the Deutsche Forschungsgemeinschaft as part of a collaborative research group on 'Discrimination and Tolerance in Intergroup Relations' (DFG, FOR 481/1-1,2).

### REFERENCES

- Allan, L. G. (1980). A note on measurement of contingency between two binary variables in judgment tasks. *Bulletin of the Psychonomic Society*, 15, 147–149.
- Batchelder, W. H., & Riefer, D. M. (1999). Theoretical and empirical review of multinomial process tree modeling. *Psychonomic Bulletin & Review*, 6, 57–86.
- Bayen, U. J., Murnane, K., & Erdfelder, E. (1996). Source discrimination, item detection, and multinomial models of source monitoring. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 22, 197–215.
- Bayen, U. J., Nakamura, G. V., Dupuis, S. E., & Yang, C.-L. (2000). The use of schematic knowledge about sources in source monitoring. *Memory & Cognition*, 28, 480–500.
- Berndsen, M., & Spears, R. (1997). Reinterpreting illusory correlation: From biased covariation to meaningful categorisation. *Swiss Journal of Psychology*, 56, 127–138.
- Fiedler, K. (1991). The tricky nature of skewed frequency tables: An information loss account of distinctiveness-based illusory correlations. *Journal of Personality and Social Psychology*, 60, 24–36.
- Fiedler, K. (1996). Explaining and simulating judgment biases as an aggregation phenomenon in probabilistic, multiple-cue environments. *Psychological Review*, 103, 193–214.
- Fiedler, K. (2000). Illusory correlations: A simple associative algorithm provides a convergent account of seemingly divergent paradigms. *Review of General Psychology*, 4, 25–58.
- Fiedler, K., Russer, S., & Gramm, K. (1993). Illusory correlations and memory performance. *Journal of Experimental Social Psychology*, 29, 111–136.
- Fiedler, K., Walther, E., Freytag, P., & Nickel, S. (2003). Inductive reasoning and judgment interference: Experiments on Simpson's paradox. *Personality and Social Psychology Bulletin*, 29, 14–27.
- Fiedler, K., Walther, E., Freytag, P., & Stryczek, E. (2002). Playing mating games in foreign cultures: A conceptual framework and an experimental paradigm for trivariate statistical inference. *Journal of Experimental Social Psychology*, 38, 14–30.
- Hamilton, D. L. (1981). Illusory correlation as a basis for stereotyping. In D. L. Hamilton (Ed.), *Cognitive processes in stereotyping and inter-group behavior* (pp. 115–144). Hillsdale, NJ: Lawrence Erlbaum.
- Hamilton, D. L., Dugan, P. M., & Trolie, T. K. (1985). The formation of stereotypic beliefs: Further evidence for distinctiveness-based illusory correlations. *Journal of Personality and Social Psychology*, 48, 5–17.



- Hamilton, D. L., & Gifford, R. K. (1976). Illusory correlation in interpersonal perception: A cognitive basis of stereotypic judgments. *Journal of Experimental Social Psychology*, 12, 392–407.
- Haslam, S. A., McGarty, C., & Brown, P. M. (1996). The search for differentiated meaning is a precursor to illusory correlation. *Personality and Social Psychology Bulletin*, 22, 611–619.
- Hintzman, D. L. (1986). 'Schema abstraction' in a multiple-trace memory model. *Psychological Review*, 93, 411–428.
- Johnson, C., & Mullen, B. (1994). Evidence for the accessibility of paired distinctiveness in distinctiveness-based illusory correlation in stereotyping. *Personality and Social Psychology Bulletin*, 20, 65–70.
- Jones, R. A., Scott, J., Solernou, J., Noble, A., Fiala, J., & Miller, K. (1977). Availability and formation of stereotypes. *Perceptual and Motor Skills*, 44, 631–638.
- Jurica, P. J., & Shimamura, A. P. (1999). Monitoring item and source information: Evidence for a negative generation effect in source memory. *Memory & Cognition*, 27, 648–656.
- Klauer, K. C., & Meiser, T. (2000). A source-monitoring analysis of illusory correlations. *Personality and Social Psychology Bulletin*, 26, 1074–1093.
- McClelland, J. L., & Rumelhart, D. E. (1985). Distributed memory and the representation of general and specific information. *Journal of Experimental Psychology: General*, 114, 159–188.
- McConnell, A. R., Sherman, S. J., & Hamilton, D. L. (1994). Illusory correlation in the perception of groups: An extension of the distinctiveness-based account. *Journal of Personality and Social Psychology*, 67, 414–429.
- McGarty, C., & de la Haye, A.-M. (1997). Stereotype formation: Beyond illusory correlation. In R. Spears, P. J. Oakes, N. Ellemers, & S. A. Haslam (Eds.), *The social psychology of stereotyping and group life* (pp. 144–170). Oxford: Blackwell.
- McGarty, C., Haslam, S. A., Turner, J. C., & Oakes, P. J. (1993). Illusory correlation as accentuation of actual intercategory difference: Evidence for the effect with minimal stimulus information. *European Journal of Social Psychology*, 23, 391–410.
- Meiser, T. (2003). Effects of processing strategy on episodic memory and contingency learning in group stereotype formation. *Social Cognition*, 21, 121–156.
- Meiser, T. (2005). A hierarchy of multinomial models for multidimensional source monitoring. *Methodology*, 1, 2–17.
- Meiser, T., & Bröder, A. (2002). Memory for multidimensional source information. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 28, 116–137.
- Meiser, T., & Hewstone, M. (2001). Crossed categorization effects on the formation of illusory correlations. *European Journal of Social Psychology*, 31, 443–466.
- Meiser, T., & Hewstone, M. (2004). Cognitive processes in stereotype formation: The role of correct contingency learning for biased group judgments. *Journal of Personality and Social Psychology*, 87, 599–614.
- Mullen, B., & Johnson, C. (1990). Distinctiveness-based illusory correlations and stereotyping: A meta-analytic integration. *British Journal of Social Psychology*, 29, 11–28.
- Pratto, F., & John, O. P. (1991). Automatic vigilance: The attention-grabbing power of negative social information. *Journal of Personality and Social Psychology*, 61, 380–391.
- Rosenberg, S., Nelson, C., & Vivekananthan, P. S. (1968). A multidimensional approach to the structure of personality impressions. *Journal of Personality and Social Psychology*, 9, 283–294.
- Rumelhart, D. E., Smolensky, P., McClelland, J. L., & Hinton, G. E. (1986). Schemata and sequential thought processes in PDP Models. In J. L. McClelland, D. E. Rumelhart, and the PDP Research Group (Eds.), *Parallel distributed processing. Explorations in the microstructure of cognition: Psychological and biological models* (Vol. 2, pp. 8–57). Cambridge, MA: MIT Press.
- Sanbonmatsu, D. M., Shavitt, S., & Gibson, B. D. (1994). Salience, set size, and illusory correlation: Making moderate assumptions about extreme targets. *Journal of Personality and Social Psychology*, 66, 1020–1033.
- Schaller, M. (1992a). In-group favoritism and statistical reasoning in social inference: Implications for formation and maintenance of group stereotypes. *Journal of Personality and Social Psychology*, 63, 61–74.
- Schaller, M. (1992b). Sample size, aggregation, and statistical reasoning in social inference. *Journal of Experimental Social Psychology*, 28, 65–85.
- Schaller, M. (1994). The role of statistical reasoning in the formation, preservation and prevention of group stereotypes. *British Journal of Social Psychology*, 33, 47–61.
- Schaller, M., Asp, C. H., Rosell, M. C., & Heim, S. J. (1996). Training in statistical reasoning inhibits the formation of erroneous group stereotypes. *Personality and Social Psychology Bulletin*, 22, 829–844.
- Schaller, M., & O'Brien, M. (1992). 'Intuitive analysis of covariance' and group stereotype formation. *Personality and Social Psychology Bulletin*, 18, 776–785.

- Simpson, E. H. (1951). The interpretation of interaction in contingency tables. *Journal of the Royal Statistical Society, Series B*, 13, 238–241.
- Skowronski, J. J., & Carlston, D. E. (1989). Negativity and extremity biases in impression formation: A review of explanations. *Psychological Bulletin*, 105, 131–142.
- Smith, E. R. (1991). Illusory correlation in a simulated exemplar-based memory. *Journal of Experimental Social Psychology*, 27, 107–123.
- Smith, E. R. (1996). What do connectionism and social psychology offer each other? *Journal of Personality and Social Psychology*, 70, 893–912.
- Smith, E. R., & DeCoster, J. (1998). Knowledge acquisition, accessibility, and use in person perception and stereotyping: Simulation with a recurrent connectionist network. *Journal of Personality and Social Psychology*, 74, 21–35.
- Snodgrass, J. G., & Corwin, J. (1988). Pragmatics of measuring recognition memory: Applications to dementia and amnesia. *Journal of Experimental Psychology: General*, 117, 34–50.
- Stroessner, S. J., Hamilton, D. L., & Mackie, D. M. (1992). Affect and stereotyping: The effect of induced mood on distinctiveness-based illusory correlations. *Journal of Personality and Social Psychology*, 62, 564–576.
- Stroessner, S. J., & Plaks, J. E. (2001). Illusory correlation and stereotype formation: Tracing the arc of research over a quarter century. In G. B. Moskowitz (Ed.), *Cognitive social psychology: The Princeton symposium on the legacy and future of social cognition* (pp. 247–259). Mahwah, NJ: Lawrence Erlbaum.
- van Rooy, D., van Overwalle, F., Vanhoomissen, T., Labiouse, C., & French, R. (2003). A recurrent connectionist model of group biases. *Psychological Review*, 110, 536–563.
- Wegener, I., & Klauer, K. C. (2004). Inter-category versus intra-category fit: When social categories match social context. *European Journal of Social Psychology*, 34, 567–593.